Imperial College London

PROBABILISTIC INFERENCE (CO-493)

Gaussian Processes

Gaussian Processes

The intention of this coursework is for you to get a better understanding of Gaussian processes by implementing Gaussian process regression.

Provided to you are **gp_assignment.py**, a skeleton file in which you will provide solutions, and **boston_housing.txt**, a file containing a dataset on house pricing that we will be using to test your solutions. Information about the dataset can be found here: `https://www.cs.toronto.edu/~delve/data/boston/bostonDetail.html`. You are given the function **loadData** which returns the Boston Housing dataset partitioned into a training set, $\mathcal{D}_{\text{train}} = \{X, \mathbf{y}\}$, and a test set, $\mathcal{D}_{\text{test}} = \{X_*, \mathbf{y}_*\}$.

If at any point you are experiencing trouble with numerical stability, the mathematical appendix in the book by Rasmussen & Williams may give some helpful pointers: `http://www.gaussianprocess.org/gpml/`.

Submit your final version to CATe via the LabTS system.

Generally, we consider the regression setting

$$y = f(\boldsymbol{x}) + \epsilon, \quad \epsilon \sim \mathcal{N}\left(0, \sigma_n^2\right).$$

We place a GP prior on f with mean function $m \equiv 0$ and covariance function k.

Task 1: 10 marks

Complete the definition of the function **multivariateGaussianSample**. This function takes a mean vector, μ, and covariance matrix, Σ, and returns a sample drawn from $\mathcal{N}(\mu, \Sigma)$. This is useful if you want to visualize draws from a GP prior or posterior.

Task 2: 20 marks

Complete the definition of the function **covMatrix**. In this part, we are considering an additive kernel/covariance function: a linear kernel plus the squared exponential (Gaussian/radial-basis-function) kernel with a White-Noise kernel to account for the Gaussian likelihood (noise model):

$$k(\mathbf{x}_p, \mathbf{x}_q) = k_{\text{Linear}}(\mathbf{x}_p, \mathbf{x}_q) + k_{\text{RBF}}(\mathbf{x}_p, \mathbf{x}_q) \tag{1}$$

$$k_{\text{Linear}}(\mathbf{x}_p, \mathbf{x}_q) = \sigma_b^2 + \sigma_v^2 \, \mathbf{x}_p \cdot \mathbf{x}_q$$

$$k_{\text{RBF}}(\mathbf{x}_p, \mathbf{x}_q) = \sigma_f^2 \exp\left(-\frac{1}{2\ell^2}\|\mathbf{x}_p - \mathbf{x}_q\|^2\right) + \sigma_n^2 \delta_{pq} \qquad \delta_{pq} = \begin{cases} 1 & \text{if } p = q \\ 0 & \text{otherwise} \end{cases}$$

We included the contribution of the Gaussian likelihood in the kernel definition ($\sigma_n^2 \delta_{pq}$) as this will simplify the implementation of the model.

For reasons explained in Task 5, we will consider the parameters $\sigma_b^2, \sigma_v^2, \sigma_f^2, \ell, \sigma_n^2$ directly **instead** of the log-parameters $\ln \sigma_b, \ln \sigma_v, \ln \sigma_f, \ln \ell, \ln \sigma_n$ given by the identities in (**??**). This is why the class **LinearPlusRBF** is initialized using log-parameters, but computes the value of the parameters via the identities.

The function **covMatrix** should return the kernel matrix K, where $K = K(A, A)$ is the kernel matrix computed for a set of points, A, using equation (**??**).

Task 3: 25 marks

Complete the definition of the function **predict**, which takes a set of test points X_* and computes the posterior mean, $\bar{\mathbf{f}}_*$, and covariance, $\mathrm{cov}(\mathbf{f}_*)$, of the GP regression for X_*.

Task 4: 5 marks

Complete the definition of the function **logMarginalLikelihood**, which computes the *negative log marginal likelihood* of the training set. Note: our optimizer, provided for you in the function **optimize**, minimizes the target function, so **please return the negative log marginal likelihood**:

$$-\log p(\mathbf{y}|X) = \frac{1}{2}\mathbf{y}^\top K^{-1}\mathbf{y} + \frac{1}{2}\log|K| + \frac{n}{2}\log 2\pi \tag{2}$$

Task 5: 25 marks

Complete the definition of the function **gradLogMarginalLikelihood**, which computes the gradients of the negative log marginal likelihood you found in Task 4. The function **optimize** will minimize the negative log-marginal likelihood on the training set using these gradients via the BFGS algorithm.
Note: we can optimize the parameters of the GP using constraints $\sigma_b^2, \sigma_v^2, \sigma_f^2, \ell, \sigma_n^2 > 0$, but a simpler method would be to solve the unconstrained optimization problem for the log parameters using the identities:

$$\begin{aligned}
\sigma_b^2 &= e^{2\ln \sigma_b} \\
\sigma_v^2 &= e^{2\ln \sigma_v} \\
\sigma_f^2 &= e^{2\ln \sigma_f} \\
\ell &= e^{\ln \ell} \\
\sigma_n^2 &= e^{2\ln \sigma_n}
\end{aligned} \tag{3}$$

Optimization is accomplished by replacing each instance of $\sigma_b^2, \sigma_v^2, \sigma_f^2, \ell, \sigma_n^2$ in equation (**??**) with the corresponding identity in (**??**) and differentiating the rewritten negative log-marginal likelihood with respect to the log parameters $\ln \sigma_b, \ln \sigma_v, \ln \sigma_f, \ln \ell, \ln \sigma_n$.

Task 6: 5 marks

Using **optimize** and initial parameter values $\{\sigma_b^2 = \sigma_v^2 = \sigma_f^2 = 1.0, \ell = 0.1, \sigma_n^2 = 0.5\}$, find the optimal parameters for the GP regression. Note the initial log parameter values are $\{\log \sigma_b = \log \sigma_v = \log \sigma_f = 0.5 \log(1.0), \log \ell = \log(0.1), \log \sigma_n = 0.5 \log(0.5)\}$.

Task 7: 10 marks

Complete the definitions of the test statistics functions **mse** and **msll** and compute the MSE and MSLL for the test set using your trained GP regression.
The function **mse** computes the *mean squared error* on the test set $\{X_*, \mathbf{y}_*\}$ using the observed values $\mathbf{y}_*$ and the predictions, $\bar{\mathbf{f}}_*$, for the test input values, X_*.

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^{n} \left(y_*^{(i)} - \bar{f}(\mathbf{x}_*^{(i)}) \right)^2 \tag{4}$$

The function **msll** computes the *mean standardized log loss* on the test set $\{X_*, \mathbf{y}_*\}$ using the observed values $\mathbf{y}_*$ and the predictions, $\bar{\mathbf{f}}_*$, for the test input values, X_*, and $\text{cov}(\mathbf{f}_*)$, the covariance of the predictive distribution of the test data.

$$\text{MSLL} = \frac{1}{n} \sum_{i=1}^{n} -\log p(y_*^{(i)} | \mathcal{D}_{\text{train}}, \mathbf{x}_*^{(i)}) = \frac{1}{n} \sum_{i=1}^{n} \frac{1}{2} \log(2\pi \sigma^2(\mathbf{x}_*^{(i)})) + \frac{\left(y_*^{(i)} - \bar{f}(\mathbf{x}_*^{(i)}) \right)^2}{2\sigma^2(\mathbf{x}_*^{(i)})} \tag{5}$$

$\sigma^2(\mathbf{x}_*^{(i)})$ is the predictive variance given by $\sigma^2(\mathbf{x}_*^{(i)}) = \mathbb{V}(f_*^{(i)}) + \sigma_n^2$ and $\mathbb{V}(f_*^{(i)})$ denotes the predictive variance of the function value for test case i.